

Virtual Storage Manager for Ceph

Version: 2.0.0.149

Source: 2015-07-31

Keywords: Ceph, Openstack, Virtual Storage Management

Supported Combo:

```
OS:           Ubuntu 14.04.2
Ceph:         Firefly/Giant/Hammer
OpenStack:    Havana/Icehouse/Juno
```

(Other combos might also be working, but we didn't try yet.)

Preparation

Before you get ready to install VSM, you should prepare your environment. **VSM CANNOT manage Ceph Cluster not created by it.** The sections here are helpful for understanding the deployment concepts.

Note: For a Ceph cluster created and managed by VSM you need to prepare at least three storage nodes plus a VSM controller node. VSM requires a minimum of three Ceph storage nodes (physical or virtual) before it will create a Ceph cluster.

Roles

There are two roles for the nodes (servers) on your VSM created Ceph cluster.

Controller Node

The controller node is used to run mariadb, rabbitmq, web ui services for the VSM cluster.

Agent Node (a.k.a Storage Node)

The agent node is used to run the vsm-agent service which manages the Ceph and physical storage resources. These nodes are the Ceph storage and monitor nodes.

Network

There are three kinds of networks defined in VSM, and the three networks can all be the same network or separate networks or subnets. VSM does not support split subnets - e.g. two or more different subnets that together make up the management network, or the ceph public network. or the ceph public network.

Management Network

Management Network is used to manage the VSM cluster, and interchanges VSM mangement data between vsm controller and agents.

Ceph Public Network

Ceph Public Network is used to serve IO operations between ceph nodes and clients.

Ceph Cluster Network

Ceph Cluster Network is used to interchange data between ceph nodes like Monitors and OSDs for replication and rebalancing.

Recommendations

- Controller node should have connectivity to:

```
Management Network
```

- Agent Node should have connectivity to:

```
Management Network
Ceph Public Network
Ceph Cluster Network
```

Sample 1

- **Controller node** contains the networks listed below:

```
192.168.123.0/24
```

- **Storage node** contains networks below:

```
192.168.123.0/24
192.168.124.0/24
192.168.125.0/24
```

Then we may assign these networks as below:

```
> Management network: 192.168.123.0/24
> Ceph public netwok: 192.168.124.0/24
> Ceph cluster network: 192.168.125.0/24
```

The configuration for VSM in the `cluster.manifest` file should be:

```
> [management_addr]
> 192.168.123.0/24
>
> [ceph_public_addr]
> 192.168.124.0/24
>
> [ceph_cluster_addr]
> 192.168.125.0/24
```

Refer [cluster.manifest](#) for details.

Sample 2

But how about when all the nodes just have two NICs. Such as a controller node and storage node having:

```
> 192.168.123.0/24
> 192.168.124.0/24
```

We can assign these two networks as below:

```
> Management network: 192.168.123.0/24
> Ceph public network: 192.168.124.0/24
> Ceph cluster network: 192.168.123.0/24
```

The configuration for VSM in `cluster.manifest` file would then be:

```
> [management_addr]
> 192.168.123.0/24
>
> [ceph_public_addr]
> 192.168.124.0/24
>
> [ceph_cluster_addr]
> 192.168.123.0/24
```

Sample 3

It's quite common to have just one NIC in demo environment, then all nodes just have:

```
> 192.168.123.0/24
```

We may assign this network as below:

```
> Management network: 192.168.123.0/24
> Ceph public network: 192.168.123.0/24
> Ceph cluster network: 192.168.123.0/24
```

So all of the three VSM networks use the same subnet, The configurations in `cluster.manifest` file would then be:

```
> [management_addr]
> 192.168.123.0/24
>
> [ceph_public_addr]
> 192.168.123.0/24
>
> [ceph_cluster_addr]
> 192.168.123.0/24
```

Automatic Deployment

Starting with VSM 1.1, an automatic deployment tool is provided which can simplify the deployment. This tool is still in development, so your feedback and JIRA reports of any problems are very welcome.

This section will describe how to use the tool to conduct automation.

1. Firstly, a VSM binary release package should be acquired. It may be downloaded from binary repository, or built from source (see [Build VSM](#)). Then unpack the release package, the folder structure looks as following:

```
.
├── CHANGELOG
├── hostrc
├── INSTALL.md
├── install.sh
├── uninstall.sh
├── LICENSE
├── manifest
│   ├── cluster.manifest.sample
│   └── server.manifest.sample
├── NOTICE
├── README
└── vsmrepo
    ├── python-vsmclient_2.0.0-123_amd64.deb
    ├── Packages.gz
    ├── vsm_2.0.0-123_amd64.deb
    ├── vsm-dashboard-2.0.0-123_amd64.deb
    └── vsm-deploy-2.0.0-123_amd64.deb
```

2. Changing the `hostrc` file, set the `AGENTADDRESSLIST` and the `CONTROLLER_ADDRESS`, the ip addresses in `AGENTADDRESSLIST` is delimited by space, and all ip addresses are used in management subnet. e.g.:

```
AGENT_ADDRESS_LIST="192.168.123.21 192.168.123.22 192.168.123.23"
CONTROLLER_ADDRESS="192.168.123.10"
```

It's OK to use host name instead of ip addresses here.

3. VSM will sync `/etc/hosts` file from the controller node, make sure your controller node's `/etc/hosts` file follows below rules:

Lines with `localhost`, `127.0.0.1` and `::1` should not contains the actual hostname.

- Under the *manifest* folder, you should create the folders named by the management ip of the controller and storage nodes, and then the structure looks as follows:

```
.
├── 192.168.123.10
├── 192.168.123.21
├── 192.168.123.22
├── 192.168.123.23
├── cluster.manifest.sample
└── server.manifest.sample
```

- Copy the *cluster.manifest.sample* to the folder named by the management ip of controller node, then change the filename to *cluster.manifest* and edit it as required, refer [cluster.manifest](#) for details.
- Copy the *server.manifest.sample* to the folders named by the management ip of storage nodes, then change the filename to *server.manifest* and edit it as required, refer [server.manifest](#) for details.
- Finally, the manifest folder structure looks as follows:

```
.
├── 192.168.123.10
│   └── cluster.manifest
├── 192.168.123.21
│   └── server.manifest
├── 192.168.123.22
│   └── server.manifest
├── 192.168.123.23
│   └── server.manifest
├── cluster.manifest.sample
└── server.manifest.sample
```

- If you want to upgrade vsm binary packages only, one approach is to build release package separately (see [Build Packages](#)). The generated binary packages will be in *vsmrepo* folder after unpack the release package, then you can execute below command to install binary package:

```
dpkg -i <package>
```

- Now we are ready to start the automatic procedure by executing this command line:

```
sudo ./install.sh -u ubuntu -v <version>
```

where *version* is the vsm version like 1.1, 2.0.

- If execution is blocked at any point, please try to enter "y" and move ahead.
- If all goes well, you can then [login to the VSM Web UI](#).

VSM Web UI

- Access <https://vsm controller IP/dashboard/vsm>.(for example <https://192.168.123.10/dashboard/vsm>)
- User name: admin, and password can be obtained from: `/etc/vsmdeploy/deployrc` in the ADMIN_PASSWORD field:

```
./get_pass.sh
```

- Then you can switch to the `Cluster Management` item, then `Create Cluster` panel, and push the create cluster button to create a ceph cluster. At this point please refer to the VSM Manual, which is located at <https://01.org/virtual-storage-manager>

Uninstall

There are a few cases where you may expect to uninstall VSM, e.g, you expect to reinstall it with different configurations, you feel VSM doesn't

work as you expected. You could take below steps to do the removal:

1. Go to the VSM folder where you start the installation procedure.
2. Make sure the `hostrc` file is there, and the ip addresses for controller node and agent nodes are correctly set. Normally, if you correctly installed VSM, you should have already correctly set the file.
3. Execute below command:

```
./uninstall.sh
```

Reference

Build Packages

There are two approaches to get a VSM release package, a direct way is to download release package from [github](#), or you can build release package from source code as following:

```
> ./buildvsm.sh -v <version>
```

where *version* is the vsm version like 1.1, 2.0. A release package named like *2.0.0-123.tar.gz* will be generated in *release* folder if all execute well.

cluster.manifest

The cluster.manifest file is under manifest// folder, the three subnets must be modified according to Ceph cluster network topology.

subnets

1. Modify the three IP addresses according to your environment. `management_addr` is used by VSM to communicate with different services, such as using rabbitmq to transfer messages, rpc.call/rpc.cast etc. `ceph_public_addr` is a public (front-side) network address. `ceph_cluster_addr` is a cluster (back-side) network address.

Also, make sure the netmask is correctly set. In this sample, *netmask=24* is fine, but with AWS instances, normally, *netmask=16* are required.

```
[management_addr]
192.168.123.0/24

[ceph_public_addr]
192.168.124.0/24

[ceph_cluster_addr]
192.168.125.0/24
```

Here is a complete list of all settings for cluster.manifest:

- **[storage_class]**

In this section, you can put you planned storage class name. One line for one class name, only names with numbers, alphabetic and underscore can be used for class name.

- **[storage_group]**

In this section, you can put your storage group definition, in the below format, here [] is not needed. Only numbers, alphabetic and underscore can be used for any of them.

```
[storage group name] [user friendly storage group name] [storage class name]
```

- **[cluster]**

In this section, you can put your cluster name. Only numbers, alphabetic and underscore can be used.

- **[file_system]**

You can use the file system which ceph can support here. The default value is xfs

- **[zone]**

In this section, you can add zone name under the section.

- format:

```
[zone-name]
```

- comments:

1. Only numbers, alphabetic and underscore can be used for zone name.
2. By default, this section is disabled, in this case, a default zone called *zone_one* will be used.

- example:

```
zone1
```

- **[management_addr]**

- **[ceph_public_addr]**

- **[ceph_cluster_addr]**

Those 3 sections will define the three subnets.

- **[settings]**

In the section, you can set values for these settings for ceph and VSM.

```
storage_group_near_full_threshold 65
storage_group_full_threshold 85
ceph_near_full_threshold 75
ceph_full_threshold 90
pg_count_factor 100
heartbeat_interval 5
osd_heartbeat_interval 10
osd_heartbeat_grace 10
disk_near_full_threshold 75
disk_full_threshold 90
osd_pool_default_size 3
```

- **[ec_profiles]**

In this section, you can define some erasure coded pool profile before you create the cluster.

- format:

```
profile-name] [path-to-plugin] [plugin-name] [pg_num value] [json format key/value]
```

- comments:

1. the key/value strings should not have spaces.

- example:

```
default_profile /usr/lib64/ceph/erasure-code jerasure 3 {"k":2,"m":1,"technique":"reed_sol_van"}
```

- **[cache_tierdefaults]**

The default settings value for create cache tier in the web UI. You can also change them while you create cache tier for pools.

```
ct_hit_set_count 1
ct_hit_set_period_s 3600
ct_target_max_mem_mb 1000000
ct_target_dirty_ratio 0.4
```

```
ct_target_full_ratio 0.8
ct_target_max_objects 1000000
ct_target_min_flush_age_m 10
ct_target_min_evict_age_m 20
```

server.manifest

The server.manifest file is under manifest// folder, below settings must be modified based on your environment.

- **[vsm_controller_ip]**

Here `vsm_controller_ip` is the VSM controller's IP address under `management_addr` subnet.

- example:

```
[vsm_controller_ip]
192.168.123.10
```

- **[role]**

Delete one if you don't want this server act as this role. The default is that server will act as storage node and monitor at the same time.

- example:

```
[role]
storage
monitor
```

- **[auth_key]**

Replace the content with the key you get from controller by running the agent-token command on the controller.

DON'T MODIFY IT, the automatic deployment tool will fill this section.

- **OSD definition under each storage group**

The storage you use for your Ceph cluster must have previously been provisioned by you with a label and a partition.

For example:

```
parted /dev/sdb -- mklabel gpt
parted -a optimal /dev/sdb -- mkpart xfs 1MB 100%
```

Enter your primary and associated journal storage information in the server.manifest, remember to fill them in right storage group.

For example, change the lines below:

```
[10krpm_sas]
#format [sas_device] [journal_device]
%osd-by-path-1% %journal-by-path-1%
%osd-by-path-2% %journal-by-path-2%
%osd-by-path-3% %journal-by-path-3%
```

to be:

```
[10krpm_sas]
#format [sas_device] [journal_device]
/dev/sdb1 /dev/sdc1
/dev/sdd1 /dev/sdc2
/dev/sde1 /dev/sdf
```

Then delete the redundant lines with `%osd-by-path%`, if you have fewer disks.

We recommend though that you use disk-by-path instead for the disk paths. Use the command below to find the true by-path:

```
ls -al /dev/disk/by-path/* | grep `disk-path` | awk '{print $9,$11}'
```

For example:

```
$> ls -al /dev/disk/by-path/* | grep sdb | awk '{print $9,$11}'  
/dev/disk/by-path/pci-0000:00:0c.0-virtio-pci-virtio3 ../../sdb
```

Then replace the `/dev/sdb` with `/dev/disk/by-path/pci-0000:00:0c.0-virtio-pci-virtio3` in `/etc/manifest/server.manifest` file. Do this also for all the other disks listed in this file.

Warning: It may cause an error when you add a disk without by-path. So, If you can not find the by-path for a normal disk, you should not use it. Or if you use it to create the cluster, and the create cluster fails, please delete it from the `/etc/manifest/server.manifest` file.

After that the disk list appears like this, here the storage group name `10krpm_sas` should have already defined in `[storage_group]` section in `cluster.manifest`.

```
[10krpm_sas]  
#format [sas_device] [journal_device]  
/dev/disk/by-path/pci-0000:00:0c.0-virtio-pci-virtio3 /dev/disk/by-path/pci-0000:00:0d.0-virtio-pci-virtio4  
/dev/disk/by-path/pci-0000:00:0e.0-virtio-pci-virtio5 /dev/disk/by-path/pci-0000:00:0f.0-virtio-pci-virtio6  
/dev/disk/by-path/pci-0000:00:10.0-virtio-pci-virtio7 /dev/disk/by-path/pci-0000:00:11.0-virtio-pci-virtio8
```

Storage Groups

If you have several kinds of storage media like 10krpm SAS drives & solid state drives, and you want these disks organized into different storage groups in VSM, then you may follow the operations below. Otherwise, you may skip this step and just put all the disks into the `[10krpm_sas]` section.

You may want to add disks into other sections in the `/etc/manifest/server.manifest` file after the `[10krpm_sas]` section. Take `[ssd]` as an example:

- Add storage class in `/etc/manifest/cluster.manifest` in controller node.

```
[storageclass] ssd # add this line 10krpmsas
```

- Add storage group in `/etc/manifest/cluster.manifest` on the controller node.

```
[storagegroup] highperformancetest HighPerformanceSSDtest ssd
```

- Add disks under `/etc/manifest/server.manifest` on the storage node(s) which have SSD, such as:

```
[ssd]  
/dev/disk/by-path/pci-0000:00:0c.0-virtio-pci-virtio9 /dev/disk/by-path/pci-0000:00:0d.0-virtio-pci-virtio1  
/dev/disk/by-path/pci-0000:00:0e.0-virtio-pci-virtio11 /dev/disk/by-path/pci-0000:00:0f.0-virtio-pci-virtio14  
/dev/disk/by-path/pci-0000:00:10.0-virtio-pci-virtio23 /dev/disk/by-path/pci-0000:00:11.0-virtio-pci-virtio10
```

Trouble shooting

If you encountered any issues, below files may give you hints:

```
> /var/log/vsm/*.log  
> /var/log/httpd/*.log  
> /var/log/syslog
```

If you want to check that the vsm agent started correctly, you can look at the files in `/var/log/vsm` on each storage node. The `vsm.physical.log` should have no errors and end with the line:

```
INFO [vsm.openstack.common.rpc.common] Connected to AMQP server on <vsm-controller-IP-address>:5673
```

and you should see a similar message at the start of `vsm.agent.log` and ending with:


```
INFO [vsm.agent.manager] agent/manager.py update ceph.conf from db. OVER
```

Likewise, you should see no errors in the three log files in `/var/log/vsm` on the controller node.

Frequently Asked Questions

Q: Executing "agent-token" is hung.

A: Please check http proxy setting to make sure no `http_proxy` variable is set in the environment.

Q: "An error occurred authenticating. Please try again later." appears on the controller web ui after fresh installation.

A: Firstly, please make sure the right password is entered, the password can be obtained from `/etc/vsmdeploy/deployrc` in "ADMIN_PA

Q: keyring error on cluster creation.

A: The root cause is that the vsm controller has already updated a new token, but it is not applied on all agents.

Q: Negative update time is showing on RBD list page.

A: Before creating the ceph cluster, please make sure all ceph nodes are time synchronized via NTP.

Q: vsm-agent process causes one disk to be saturated with i/o load.

A: A known case causes i/o saturation if multiple OSDs are defined on the same physical device, which is normally used in the demo

Q: Can't replace node if ceph cluster contains only 3 nodes.

A: This is an expected safeguard. A 3 node cluster minimum is needed to meet availability requirements.

Q: Can I define the ceph version I want to install?

A: VSM can work with different Ceph releases like Firefly, Giant and Hammer. By default, it will use the ceph version provided by Ceph. If user expects to use some specific ceph version, he need add the new ceph repo into `/etc/apt/sources.list.d/ceph.list` to override the default one.

```
>
> sudo rm /etc/apt/sources.list.d/ceph.list
> echo deb http://ceph.com/debian-firefly/ $(lsb_release -sc) main | sudo tee /etc/apt/sources.list.d/ceph.list
>
```